

```
In [8]: import matplotlib.pyplot as plt
import numpy as np
import os
import PIL
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential

import pathlib

dataset_url = "https://storage.googleapis.com/download.tensorflow.org/example_images/flower_photos.tgz"
data_dir = tf.keras.utils.get_file(
    'flower_photos', origin=dataset_url, untar=True)
data_dir = pathlib.Path(data_dir)

image_count = len(list(data_dir.glob('*/*.jpg')))
print(image_count)

3670
```

```
In [9]: roses = list(data_dir.glob('roses/*'))
PIL Image.open(str(roses[0]))
```



```
In [10]: # Training split
train_ds = tf.keras.utils.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="training",
    seed=123,
    image_size=(180, 180),
    batch_size=32)
```

Found 3670 files belonging to 5 classes.
Using 2936 files for training.

```
In [11]: # Testing or Validation split
val_ds = tf.keras.utils.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="validation",
    seed=123,
    image_size=(180, 180),
    batch_size=32)
```

Found 3670 files belonging to 5 classes.
Using 734 files for validation.

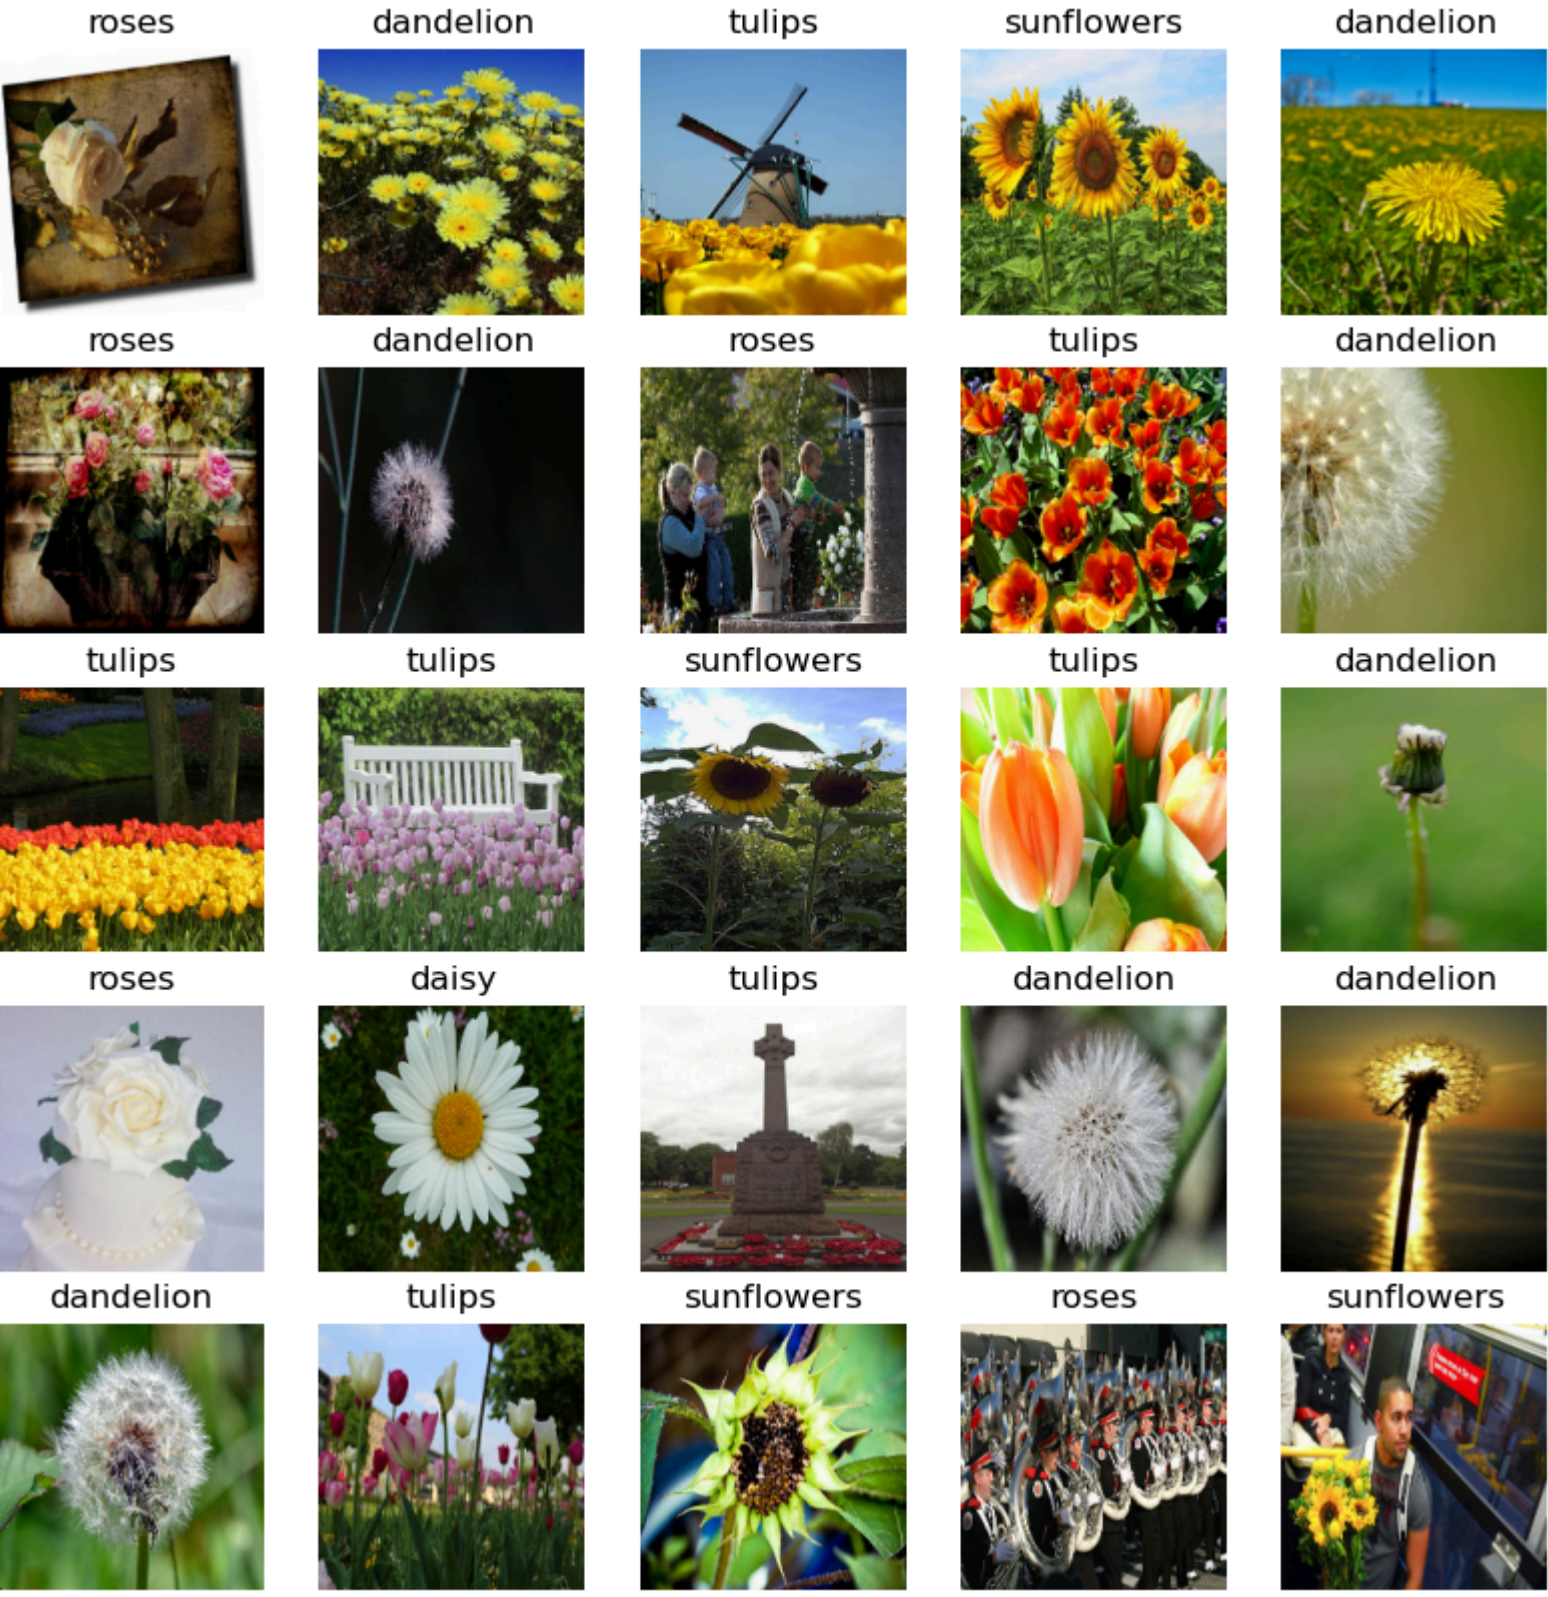
```
In [12]: class_names = train_ds.class_names
print(class_names)
```

['daisy', 'dandelion', 'roses', 'sunflowers', 'tulips']

```
In [13]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 10))

for images, labels in train_ds.take(1):
    for i in range(25):
        ax = plt.subplot(5, 5, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(class_names[labels[i]])
        plt.axis("off")
```



```
In [15]: # Define the number of classes (you need to define class_names somewhere in your code)
class_names = ['daisy', 'dandelion', 'roses', 'sunflowers', 'tulips']
num_classes = len(class_names)
```

```
# Define the model
model = Sequential([
    layers.Input(shape=(180, 180, 3)),
    layers.Rescaling(1./255),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)
])

# Print the model summary
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
rescaling_1 (Rescaling)	(None, 180, 180, 3)	0
conv2d_3 (Conv2D)	(None, 180, 180, 16)	448
max_pooling2d_3 (MaxPooling2D)	(None, 90, 90, 16)	0
conv2d_4 (Conv2D)	(None, 90, 90, 32)	4,640
max_pooling2d_4 (MaxPooling2D)	(None, 45, 45, 32)	0
conv2d_5 (Conv2D)	(None, 45, 45, 64)	18,496
max_pooling2d_5 (MaxPooling2D)	(None, 22, 22, 64)	0
flatten_1 (Flatten)	(None, 38976)	0
dense_2 (Dense)	(None, 128)	3,965,956
dense_3 (Dense)	(None, 5)	645

Total params: 3,989,285 (15.22 MB)
Trainable params: 3,989,285 (15.22 MB)
Non-trainable params: 0 (0.00 B)

```
In [23]: # Compile the model
model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])
```

```
# Train the model
epochs = 10
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs
)

# Evaluate the model on the validation set
val_loss, val_acc = model.evaluate(val_ds)

# Predict on the validation set
y_pred = np.argmax(model.predict(val_ds), axis=-1)

# Get true labels
y_true = np.concatenate([y for x, y in val_ds], axis=0)

# Generate classification report
from sklearn.metrics import classification_report

report = classification_report(y_true, y_pred, target_names=class_names)
print(report)

Epoch 1/10
92/92 144ms/step - accuracy: 0.9743 - loss: 0.0676 - val_accuracy: 0.6458 - val_loss: 1.9442
Epoch 2/10
92/92 146ms/step - accuracy: 0.9795 - loss: 0.0626 - val_accuracy: 0.6676 - val_loss: 1.8916
Epoch 3/10
92/92 189ms/step - accuracy: 0.9957 - loss: 0.0173 - val_accuracy: 0.6676 - val_loss: 2.1380
Epoch 4/10
92/92 234ms/step - accuracy: 1.0000 - loss: 0.0035 - val_accuracy: 0.6649 - val_loss: 2.2106
Epoch 5/10
92/92 237ms/step - accuracy: 0.9999 - loss: 0.0024 - val_accuracy: 0.6580 - val_loss: 2.6211
Epoch 6/10
92/92 207ms/step - accuracy: 0.9673 - loss: 0.1177 - val_accuracy: 0.6553 - val_loss: 1.8801
Epoch 7/10
92/92 209ms/step - accuracy: 0.9907 - loss: 0.0273 - val_accuracy: 0.6335 - val_loss: 2.1049
Epoch 8/10
92/92 199ms/step - accuracy: 0.9921 - loss: 0.0293 - val_accuracy: 0.6199 - val_loss: 2.5368
Epoch 9/10
92/92 230ms/step - accuracy: 0.9743 - loss: 0.0843 - val_accuracy: 0.6499 - val_loss: 2.0772
Epoch 10/10
92/92 224ms/step - accuracy: 0.9983 - loss: 0.0051 - val_accuracy: 0.6526 - val_loss: 2.1659
23/23 2s 99ms/step - accuracy: 0.6340 - loss: 2.1946

precision    recall  f1-score   support

daisy       0.19      0.19      0.19      129
dandelion   0.25      0.28      0.27      176
roses       0.19      0.18      0.19      129
sunflowers  0.16      0.14      0.15      152
tulips      0.22      0.22      0.22      157

accuracy          0.21      734
macro avg         0.20      0.20      0.20      734
weighted avg      0.21      0.21      0.21      734
```


```
In [22]: #Accuracy
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

#Loss
loss = history.history['loss']
val_loss = history.history['val_loss']

#epochs
epochs_range = range(epochs)

#Plotting graphs
plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



```
In [24]: # Compile the model
model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])
```

```
# Train the model
epochs = 20
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs
)

# Evaluate the model on the validation set
val_loss, val_acc = model.evaluate(val_ds)

# Predict on the validation set
y_pred = np.argmax(model.predict(val_ds), axis=-1)

# Get true labels
y_true = np.concatenate([y for x, y in val_ds], axis=0)

# Generate classification report
from sklearn.metrics import classification_report

report = classification_report(y_true, y_pred, target_names=class_names)
print(report)

Epoch 1/20
92/92 167ms/step - accuracy: 0.9977 - loss: 0.0115 - val_accuracy: 0.6689 - val_loss: 2.5806
Epoch 2/20
92/92 123ms/step - accuracy: 0.9925 - loss: 0.0324 - val_accuracy: 0.6580 - val_loss: 2.4683
Epoch 3/20
92/92 143ms/step - accuracy: 0.9952 - loss: 0.0234 - val_accuracy: 0.6662 - val_loss: 2.7637
Epoch 4/20
92/92 161ms/step - accuracy: 0.9921 - loss: 0.0273 - val_accuracy: 0.6267 - val_loss: 2.8589
Epoch 5/20
92/92 164ms/step - accuracy: 0.9851 - loss: 0.0471 - val_accuracy: 0.6390 - val_loss: 2.6442
Epoch 6/20
92/92 166ms/step - accuracy: 0.9913 - loss: 0.0326 - val_accuracy: 0.6622 - val_loss: 2.6565
Epoch 7/20
92/92 146ms/step - accuracy: 0.9884 - loss: 0.0317 - val_accuracy: 0.6512 - val_loss: 2.5424
Epoch 8/20
92/92 145ms/step - accuracy: 0.9982 - loss: 0.0077 - val_accuracy: 0.6512 - val_loss: 2.0663
Epoch 9/20
92/92 149ms/step - accuracy: 0.9991 - loss: 0.0034 - val_accuracy: 0.6567 - val_loss: 2.6091
Epoch 10/20
92/92 144ms/step - accuracy: 0.9991 - loss: 3.5889e-04 - val_accuracy: 0.6567 - val_loss: 2.7670
Epoch 11/20
92/92 153ms/step - accuracy: 1.0000 - loss: 2.1926e-04 - val_accuracy: 0.6580 - val_loss: 2.7617
Epoch 12/20
92/92 177ms/step - accuracy: 1.0000 - loss: 1.6937e-04 - val_accuracy: 0.6580 - val_loss: 2.8085
Epoch 13/20
92/92 205ms/step - accuracy: 1.0000 - loss: 1.2953e-04 - val_accuracy: 0.6580 - val_loss: 2.8486
Epoch 14/20
92/92 173ms/step - accuracy: 1.0000 - loss: 1.0320e-04 - val_accuracy: 0.6580 - val_loss: 2.8842
Epoch 15/20
92/92 181ms/step - accuracy: 1.0000 - loss: 8.7264e-05 - val_accuracy: 0.6594 - val_loss: 2.9161
Epoch 16/20
92/92 172ms/step - accuracy: 1.0000 - loss: 6.8263e-05 - val_accuracy: 0.6594 - val_loss: 2.9448
Epoch 17/20
92/92 169ms/step - accuracy: 1.0000 - loss: 6.4049e-05 - val_accuracy: 0.6580 - val_loss: 2.9719
Epoch 18/20
92/92 162ms/step - accuracy: 1.0000 - loss: 5.2259e-05 - val_accuracy: 0.6580 - val_loss: 2.9978
Epoch 19/20
92/92 160ms/step - accuracy: 1.0000 - loss: 4.7733e-05 - val_accuracy: 0.6580 - val_loss: 3.0218
Epoch 20/20
92/92 166ms/step - accuracy: 1.0000 - loss: 4.3818e-05 - val_accuracy: 0.6594 - val_loss: 3.0449
23/23 2s 68ms/step - accuracy: 0.6797 - loss: 2.9890

precision    recall  f1-score   support

daisy       0.18      0.19      0.18      129
dandelion   0.24      0.26      0.25      176
roses       0.18      0.15      0.16      120
sunflowers  0.20      0.20      0.20      152
tulips      0.21      0.22      0.21      157

accuracy          0.21      734
macro avg      0.20      0.20      0.20      734
```